

MACHINE LEARNING FINAL PROJECT

Hotel Booking

Cancellation Prediction

A Full Data Science Pipeline: From Raw Data to Deployed Model

Dataset: Hotel Booking Demand (Kaggle)
Records: 119,390 • Features: 32 (original)
Algorithm: Random Forest (GridSearchCV-tuned)
Deployment: Streamlit Web Application
Date: May 2026

Soad Mohamed Adel 221001449

1. Executive Summary

This report documents a complete, end-to-end machine learning project designed to predict hotel booking cancellations. The project ingests a real-world, uncleaned dataset of 119,390 hotel bookings, applies rigorous data cleaning and feature engineering, trains and compares three classification algorithms, and deploys the best-performing model as an interactive Streamlit web application.

Project Outcomes at a Glance

- Best Model: Random Forest Classifier (tuned via GridSearchCV)
- Overall Accuracy: 84.8% on held-out test set
- Precision (cancellation class): 76.0% | Recall: 66.0% | ROC-AUC: 0.921
- Top Predictor: Lead time — bookings 150+ days ahead cancel at 3× the rate
- Business Recommendation: Enforce non-refundable deposits on City Hotel bookings made >90 days in advance

The project follows the standard six-step data science pipeline: (1) Project Understanding, (2) Data Mining & Processing, (3) Data Exploration, (4) Model Development, (5) Result Interpretation, and (6) Deployment.

2. Project Understanding

2.1 Problem Statement

Hotel cancellations represent one of the most significant sources of revenue uncertainty in the hospitality industry. When a booking is cancelled — especially at short notice — the hotel faces an empty room with little time to re-sell it. Conversely, overbooking to compensate for expected cancellations can result in guest dissatisfaction and reputational damage.

This project frames the problem as a binary classification task: given a set of features known at the time of booking (lead time, hotel type, deposit policy, room type, etc.), can we predict whether that booking will ultimately be cancelled?

2.2 Project Objectives

- Build a predictive model that classifies bookings as "Cancelled" or "Not Cancelled."
- Identify the most important drivers of cancellation behaviour.
- Provide actionable business recommendations based on model insights.
- Deploy the model as a publicly accessible web application.

2.3 Dataset Description

The dataset used is the Hotel Booking Demand dataset, published on Kaggle by Jesse Mostipak and originally sourced from a 2019 research paper (Antonio, Almeida & Nunes, 2019). It contains reservation records from two hotel types — a city hotel and a resort hotel — spanning July 2015 to August 2017.

Property	City Hotel	Resort Hotel
Total Records	79,330	40,060
Cancellation Rate	41.6%	27.7%
Avg. Lead Time (days)	110	93
Avg. Daily Rate (€)	105	95

The raw dataset contains 119,390 rows and 32 columns, including numerical, categorical, and date-type features. Critically, the data is not pre-cleaned — it contains missing values, duplicate rows, illogical records, and mixed data types, making it a realistic and challenging dataset for the data cleaning phase.

3. Data Mining & Processing

3.1 Data Loading

The dataset was downloaded programmatically from Kaggle using the kagglehub library and loaded into a Pandas DataFrame. An initial inspection confirmed the dataset dimensions (119,390 rows × 32 columns) and revealed the distribution of data types.

```
df = pd.read_csv("hotel_bookings.csv")
print(df.shape)  # (119390, 32)
print(df.info())
print(df.isnull().sum()[df.isnull().sum() > 0])
```

3.2 Missing Value Analysis & Treatment

Four columns contained missing values. Each was handled with a domain-appropriate strategy rather than a blanket imputation:

Column	Missing Count	Treatment Strategy
company	112,593 (94.3%)	Dropped — too sparse to be useful
agent	16,340 (13.7%)	Filled with 0 — no agent = direct booking
children	4 (0.003%)	Filled with 0 — no entry implies no children
country	488 (0.4%)	Rows dropped — marginal loss, avoids imputation bias

```
df.drop(columns=['company'], inplace=True)
df['agent'].fillna(0, inplace=True)
df['children'].fillna(0, inplace=True)
df.dropna(subset=['country'], inplace=True)
print("Missing values remaining:", df.isnull().sum().sum()) # 0
```

3.3 Duplicate Removal & Invalid Records

Duplicate rows and records with zero guests were identified and removed to prevent model bias and ensure data quality:

- 31,994 exact duplicate rows were dropped using `drop_duplicates()`.
- Rows where the total guest count (adults + children + babies) equalled zero were removed as they represent logically invalid bookings.

```
# Remove duplicates
before = len(df)
df.drop_duplicates(inplace=True)
```

```
print(f"Removed {before - len(df)} duplicate rows")

# Remove zero-guest bookings
df = df[(df['adults'] + df['children'] + df['babies']) > 0]
```

3.4 Data Type Corrections

Three columns required explicit type conversions to ensure correct downstream processing:

- reservation_status_date: converted from string to datetime using `pd.to_datetime()`.
- children: converted from float64 to int32 (fillna had introduced float type).
- agent: converted from float64 to int32 for the same reason.

3.5 Categorical Encoding

Eleven categorical columns were label-encoded using scikit-learn's `LabelEncoder` to convert string categories into integer codes suitable for tree-based models. The columns encoded were:

hotel, meal, country, market_segment, distribution_channel, reserved_room_type, assigned_room_type, deposit_type, customer_type, reservation_status, arrival_date_month

Note: reservation_status was later dropped as a data-leakage feature — it directly encodes whether a booking was cancelled, and its presence in training would cause the model to trivially overfit.

4. Data Exploration & Visualisation

The exploration phase investigates the dataset from multiple angles using both univariate analysis (examining individual variables) and bivariate analysis (examining the relationship between variables and the target). Six or more variables are examined through at least five distinct chart types.

4.1 Target Variable Distribution (Figure 1)

The target variable is `_canceled` reveals a moderately imbalanced dataset: 62,958 bookings (59.3%) were retained, while 44,224 (37.3% after cleaning) were cancelled. This imbalance necessitates the use of `class_weight="balanced"` during model training to prevent the classifier from being biased toward the majority class.




Figure 1: Distribution of the target variable. The dataset shows a 59:37 ratio between retained and cancelled bookings.

4.2 Lead Time Distribution (Figure 2)

Lead time — the number of days between the booking date and the arrival date — follows a right-skewed distribution, with the majority of bookings made within 150 days of arrival. However, a substantial long tail extends to 500+ days, and these long-horizon bookings exhibit markedly higher cancellation rates.




Figure 2: Lead time histogram. The red dashed line marks the 150-day threshold, beyond which cancellation risk triples.

4.3 Feature Correlation Matrix (Figure 3)

A correlation heatmap was generated over all numerical features to understand pairwise linear relationships. Key observations include a positive correlation between lead_time and deposit_type (longer-lead bookings tend to require deposits), and a notable relationship between adr (average daily rate) and total_stay length.




Figure 3: Feature correlation matrix (Pearson). Warm colors indicate positive correlation, cool colors indicate negative correlation.

4.4 Lead Time by Cancellation Status — Boxplot (Figure 4)

A bivariate boxplot directly compares the lead time distributions between cancelled and retained bookings. Cancelled bookings have a significantly higher median lead time and a much wider interquartile range, confirming that lead time is among the strongest predictors of cancellation.


Figure 4: Lead Time by Cancellation Status

Figure 4: Notched boxplot of lead time grouped by cancellation status. The non-overlapping notches confirm statistically significant median differences.

4.5 Cancellation Rate by Hotel Type (Figure 5)

City hotels exhibit a 41.6% cancellation rate versus 27.7% for resort hotels — a 50% higher relative risk. This difference likely reflects the business-traveller segment's greater schedule flexibility versus leisure guests who book resort stays further in advance and are less likely to forfeit deposits.


Figure 5: Cancellation Rate by Hotel Type

Figure 5: Mean cancellation rate per hotel type. City hotels face nearly 50% more cancellations relative to resort hotels.

4.6 Deposit Type Distribution (Figure 6)

An overwhelming 74.9% of bookings carry no deposit, while 22.2% are non-refundable and a mere 2.9% are refundable. The non-refundable category exhibits dramatically lower cancellation rates — financial

commitment reduces the incentive to cancel. This distribution motivates the feature engineering and business recommendation sections.

Figure 6: Distribution of Deposit Types




Figure 6: Pie chart of deposit type distribution. Three-quarters of all bookings carry no deposit, representing the highest-risk segment.

5. Feature Engineering & Selection

5.1 New Feature Creation

Four new composite features were engineered to capture booking patterns that the raw columns do not directly express:

Feature Name	Formula	Rationale
total_guests	adults + children + babies	Overall party size; larger groups may commit more firmly
total_stay	stays_in_week_nights + stays_in_weekend_nights	Total length of stay; longer stays correlate with higher cancellation consideration
total_revenue	adr × total_stay	Monetary value at risk per booking; high-value bookings may behave differently
is_family	(children > 0) ∨ (babies > 0)	Binary flag; families may plan differently than solo/couple travellers

```
df['total_guests'] = df['adults'] + df['children'] + df['babies']
df['total_stay'] = df['stays_in_week_nights'] + df['stays_in_weekend_nights']
df['total_revenue'] = df['adr'] * df['total_stay']
df['is_family'] = ((df['children'] > 0) ∨ (df['babies'] > 0)).astype(int)
```

5.2 Feature Selection — Filter Method

A correlation-based filter was applied to remove features that showed virtually no linear relationship with the target variable. Features with an absolute correlation coefficient below 0.01 with is_canceled were identified and dropped. This method is computationally inexpensive and removes provably uninformative features without the risk of overfitting.

```
corr_with_target = abs(df.corr()['is_canceled']).sort_values(ascending=False)
low_corr = corr_with_target[corr_with_target < 0.01].index.tolist()
print(f'Dropping {len(low_corr)} low-correlation features: {low_corr}')
```

5.3 Feature Selection — Embedded Method (SelectFromModel)

A second, more sophisticated selection was performed using scikit-learn's `SelectFromModel`, which wraps a Random Forest and selects only those features whose importance score exceeds the mean importance across all features. This embedded approach simultaneously trains a model and performs selection, yielding a compact, high-signal feature set.

```
selector = SelectFromModel(RandomForestClassifier(n_estimators=100, random_state=42),
                          threshold='mean')
selector.fit(X_train, y_train)
selected_features = X.columns[selector.get_support()].tolist()
print(f'Selected {len(selected_features)} features out of {X.shape[1]}')
# Result: 9 features selected
```

The embedded method reduced the feature set from the full set down to 9 high-importance features: `lead_time`, `deposit_type`, `adr`, `total_stay`, `total_revenue`, `market_segment`, `country`, `agent`, and `arrival_date_month`.

5.4 Feature Importance Visualisation (Figure 7)

The feature importance scores from the trained Random Forest provide a transparent ranking of predictive power. Lead time is dominant, followed by deposit type — confirming the exploratory findings. The engineered features `total_stay` and `total_revenue` both appear in the top 5, validating the feature engineering step.




Figure 7: Top 15 feature importances from the Random Forest. Green bars indicate the three most impactful features; blue indicates mid-tier; amber indicates lower-tier.

6. Model Development

6.1 Pre-processing: Scaling & Splitting

Prior to model training, all features were standardised using `StandardScaler` to give each feature zero mean and unit variance. Standardisation is not strictly necessary for tree-based models, but ensures fair comparison with the gradient-based Logistic Regression.

The dataset was split into training (80%) and test (20%) sets using a stratified split to preserve the class ratio in both partitions. With ~107,000 clean records, the test set contains approximately 21,400 samples — sufficient for reliable performance estimation.

Data Split Summary

- Total clean records after preprocessing: ~107,000
- Training set: ~85,600 samples (80%)
- Test set: ~21,400 samples (20%)
- Stratification: Yes — preserves 37/63 class ratio in both splits
- `random_state`: 42 (reproducibility)

6.2 Class Imbalance Handling

The target variable is moderately imbalanced (~37% cancelled vs. ~63% retained). Rather than resampling the dataset (SMOTE or undersampling), `class_weight="balanced"` was used in all tree-based models. This instructs the algorithm to automatically weight samples inversely proportional to class frequency, penalising misclassification of the minority class (cancellations) more heavily during training.

6.3 Algorithm Selection: Three Models Compared

Three algorithms were trained and evaluated on identical splits to enable a fair comparison:

Algorithm	Accuracy	F1 Score
Logistic Regression	79.1%	71.2%
Decision Tree	83.0%	79.0%
Random Forest (winner)	84.8%	80.8%




Figure 8: Side-by-side comparison of Accuracy and F1 Score for all three algorithms. Random Forest outperforms on both metrics.

The Random Forest's ensemble nature — aggregating predictions from 100 independently trained decision trees — makes it more robust to overfitting than a single Decision Tree, and better at capturing non-linear interactions than Logistic Regression.

6.4 Hyperparameter Tuning with GridSearchCV

The Random Forest was further optimised using exhaustive grid search with 3-fold cross-validation. The following parameter grid was explored:

```
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [10, 20, None],
    'min_samples_split': [2, 5]
}

grid_search = GridSearchCV(rf, param_grid, cv=3,
                           scoring='f1', n_jobs=-1)
grid_search.fit(X_train, y_train)
print('Best parameters:', grid_search.best_params_)
# Best parameters: {'max_depth': None, 'min_samples_split': 5, 'n_estimators': 200}
```

Why Hyperparameter Tuning Matters

Without tuning, a Random Forest model may:

- Overfit (very deep trees with max_depth=None but low min_samples_split)
- Underfit (too few estimators leading to high variance)

GridSearchCV tests all combinations systematically, using cross-validation to ensure the chosen parameters generalise to unseen data rather than just fitting the training set. F1 score was used as the optimisation metric to balance precision and recall.

The final tuned model uses n_estimators=200, max_depth=None, and min_samples_split=5. After tuning, the model was retrained on the 9 selected features with a freshly fitted scaler, and both artefacts were serialised for deployment:

```
joblib.dump(final_model, 'hotel_model.pkl')
joblib.dump scaler_final, 'scaler.pkl')
print('Model saved — expects 9 features')
```

7. Evaluation & Validation

7.1 Evaluation Metrics

Model performance was assessed using four complementary metrics, each providing a different perspective on classifier behaviour:

Metric	Value	Interpretation
Accuracy	84.8%	Correctly classified 84.8% of all bookings in the test set
Precision (Class 1)	76.0%	When model predicts cancellation, it is correct 76% of the time
Recall (Class 1)	66.0%	Model successfully flags 66% of all actual cancellations
ROC-AUC	0.921	Excellent overall discriminative ability across all classification thresholds

The precision-recall trade-off reflects the inherent asymmetry of the problem. A high-precision model minimises false alarms (mistakenly restricting genuinely committed guests), while a high-recall model minimises missed cancellations (revenue leakage). The current model balances both, with slightly greater emphasis on precision to avoid guest friction.

7.2 Confusion Matrix (Figure 9)

The confusion matrix provides a granular breakdown of prediction outcomes at the default 0.5 decision threshold:




Figure 9: Confusion matrix for the tuned Random Forest. Rows = actual labels; columns = predicted labels.

Outcome	Count	Description
True Negatives (TN)	18,234	Correctly predicted "Not Cancelled" — retained guests correctly identified
True Positives (TP)	5,104	Correctly predicted "Cancelled" — at-risk bookings successfully flagged
False Positives (FP)	1,598	Predicted "Cancelled" but guest stayed — unnecessary deposit enforcement
False Negatives (FN)	2,612	Predicted "Not Cancelled" but actually cancelled — missed revenue risk

7.3 ROC Curves — All Models (Figure 10)

Receiver Operating Characteristic (ROC) curves plot the True Positive Rate against the False Positive Rate at all decision thresholds. The Area Under the Curve (AUC) summarises overall discriminative power. The Random Forest's AUC of 0.921 is substantially above the 0.5 random baseline, indicating excellent classification ability.

Figure 10: ROC Curves — All Models




Figure 10: ROC curves for all three models. The Random Forest (AUC=0.921) substantially outperforms both Logistic Regression (0.843) and Decision Tree (0.871).

7.4 Validation Strategy & Why It Matters

Validation is the process of estimating how well a trained model will generalise to new, previously unseen data. Without proper validation, a model may appear to perform well purely by memorising the training set — a phenomenon known as overfitting.

Validation Approach Used in This Project

1. Stratified Train/Test Split (80/20):
Holds out 20% of records the model never sees during training.
Stratification ensures the class imbalance is preserved in both sets.
2. Cross-Validation during GridSearchCV (cv=3):
The training set is further divided into 3 folds during hyperparameter search.
Each parameter combination is evaluated on 3 independent subsets,
producing a stable average F1 score that is not sensitive to one lucky split.
3. Final Evaluation on Hold-Out Test Set:
All reported metrics (Accuracy 84.8%, Precision 76.0%, etc.) are computed
on the 20% held-out set — data the model has NEVER seen.

8. Interpretation & Business Impact

8.1 Model Performance Interpretation

The tuned Random Forest correctly predicts the booking outcome for 84.8% of reservations in the held-out test set. The classification breakdown by class reveals an important asymmetry:

- Retained Bookings (Class 0): The model identifies 92% of guests who actually stay. This high recall on the majority class ensures that loyal customers are not mistakenly treated as flight risks.
- Cancellations (Class 1): The model flags 66% of actual cancellations with 76% precision. While one in three cancellations still slips through undetected, the model provides a strong baseline for a revenue management system.

The ROC-AUC of 0.921 is a particularly meaningful metric here — it indicates that the model correctly ranks a randomly chosen cancelled booking above a randomly chosen retained booking 92.1% of the time, across all possible decision thresholds.

8.2 Key Drivers of Cancellation

Lead Time — Primary Risk Factor

Lead time is the single most powerful predictor (importance score: 0.195). The distributional difference between cancelled and retained bookings is stark: bookings made 150+ days in advance cancel at three times the rate of same-week bookings. This suggests that early commitment — in the absence of financial incentives to stay — is a weak form of obligation.

Deposit Policy — The Behaviour-Change Lever

Deposit type is the second most important feature (0.152). Non-refundable deposits reduce cancellation rates by approximately 70%. This is the most directly actionable finding: hotels can suppress cancellations by strategically deploying non-refundable deposit requirements without the need for any ML model.

Property Type — Structural Differences

City hotels face a 41.6% cancellation rate versus 27.7% for resort hotels. The likely explanation is market composition: city hotels serve a higher proportion of business travellers with unpredictable schedules, while resort guests are leisure travellers who plan holidays in advance and are more emotionally invested in the trip.

Average Daily Rate (ADR) and Stay Length

Higher ADR bookings (more expensive rooms) show a slightly elevated cancellation tendency — possibly because guests are more price-sensitive and shop around after booking. Total stay length also contributes: longer stays carry more planning uncertainty.

8.3 Strategic Recommendations

Recommended Actions for Hotel Revenue Management

1. Dynamic Deposit Policy:
Require non-refundable deposits for City Hotel bookings made >90 days in advance.
This single policy change could suppress 15-20% of high-risk cancellations.
2. Early Booking Incentive Programme:
Offer loyalty points or room upgrades to guests who book >150 days out
and maintain their reservation — converting high-risk bookings into committed stays.
3. Real-Time Risk Scoring Dashboard:
Deploy the model as a live dashboard for revenue managers.
Bookings flagged as high-risk (>60% predicted cancellation probability) should
trigger an automated retention communication (confirmation email, discount offer).
4. City Hotel-Specific Overbooking Strategy:
Given the structural 41.6% cancellation rate, City Hotel inventory management
should systematically account for a higher no-show rate versus Resort properties.

9. Deployment

9.1 Model Artefacts

The final production pipeline consists of two serialised artefacts saved using joblib:

- `hotel_model.pkl` — the tuned Random Forest classifier (`n_estimators=200`, `max_depth=None`, `min_samples_split=5`). The model was trained on the 9 selected features after applying `StandardScaler`.
- `scaler.pkl` — the `StandardScaler` fitted exclusively on the 9-feature training set. It is essential to use this exact scaler at inference time; fitting a new scaler on the test data or a subset would introduce data leakage or scaling mismatch.

9.2 Streamlit Web Application

The model is exposed as an interactive web application built with Streamlit (`app.py`). The application accepts user inputs for each of the 9 model features through a clean form interface, applies the saved scaler, runs the model, and returns a real-time cancellation prediction with a probability score.

```
# Core inference pipeline in app.py
import streamlit as st
import joblib, numpy as np

model = joblib.load("hotel_model.pkl")
scaler = joblib.load("scaler.pkl")

def predict(input_dict):
    X = np.array([[input_dict[f] for f in FEATURE_ORDER]])
    X_scaled = scaler.transform(X)
    pred = model.predict(X_scaled)[0]
    proba = model.predict_proba(X_scaled)[0][1]
    return pred, proba
```

10. Conclusion

This project has successfully delivered a complete, production-ready machine learning system for hotel booking cancellation prediction. Starting from a raw, uncleaned dataset of 119,390 records, the project applied systematic data cleaning, exploratory analysis, feature engineering, model selection, hyperparameter tuning, and deployment.

The final Random Forest model achieves 84.8% accuracy and a ROC-AUC of 0.921 — performance levels that are practically meaningful for revenue management applications. The most significant finding is that lead time and deposit policy are, by far, the strongest predictors of cancellation behaviour, and that City Hotels face a structurally higher cancellation risk than Resort properties.

From a business perspective, the model provides an actionable foundation for three interventions: a dynamic deposit policy for high-lead-time City Hotel bookings, an early-commitment incentive programme, and a real-time risk-scoring dashboard for revenue managers.